

Sunwoo Lee

Assistant Professor
Department of Computer Engineering
Inha University, South Korea

Tel: (+82) 032-860-7445
Email: sunwool@inha.ac.kr
<https://sites.google.com/view/sunwoolee>

Research Interest

- Large-Scale Distributed Machine Learning and Deep Learning
- Federated Learning on Heterogeneous Systems
- Applied Machine Learning (E.g., Physics + ML)

Education

Northwestern University	Evanston, IL USA
Ph.D. in Computer Engineering	2020
Advisors: Prof. Alok Choudhary and Prof. Wei-keng Liao	
Thesis: Scalable Parallelization Strategy for Large-Scale Deep Learning	

Hanyang University	Seoul, South Korea
B.S. and M.S. in Computer Engineering	2009
Advisor: Prof. Minsoo Ryu	

Employment

Inha University	Incheon, South Korea
Assistant Professor of Computer Engineering	2022 – now

University of Southern California	Los Angeles, CA USA
Postdoctoral Researcher	2020 – 2022
Advisor: Prof. Salman Avestimehr	

Samsung Electronics, Memory Solutions Lab.	Hwaseong, South Korea
System Software Researcher	2013 – 2015

Humax	Bundang, South Korea
Software Engineer (Alternative Military Service)	2009 – 2013

Research Internship

Lawrence Berkeley National Laboratory
Research Intern

Berkeley, CA USA
Jun 2020 – Aug 2020

Fermi National Accelerator Laboratory
Research Intern

Batavia, IL USA
Jul 2019 – Sep 2019

Argonne National Laboratory
Research Intern

Lemont, IL USA
Jun 2018 – Aug 2018

Teaching Experience

Inha University

- CSE4315: Machine Learning Spring 2023, 2024
- CSE1103: Objected Oriented Programming 2 Spring 2023, 2024
- CSE3209: System Programming Fall 2022, 2023
- CSE3313: Linux Programming Fall 2022, 2023

University of Southern California

- AEOP Scholarship Program in Data Science Summer 2021

Northwestern University

- CE501: Social Media Mining Spring 2020
- CE303: Advanced Digital Design Fall 2019

Honors & Awards

FL-AAAI Workshop Best Paper Award 2022

- SSFL: Tackling Label Deficiency in Federated Learning via Personalized Self-Supervision

IEEE HiPC Best Paper Finalist 2017

- Parallel Deep Convolutional Neural Network Training by Exploiting the Overlapping of Computation and Communication

Publications

1. **Sunwoo Lee**, Anit Sahu, Chaoyang He, and Salman Avestimehr, Partial Model Averaging in Federated Learning: Performance Guarantees and Benefits, *Neurocomputing* 2023

2. Yue Niu, Saurav Prakash, Souvik Kundu, **Sunwoo Lee**, and Salman Avestimehr, Overcoming Resource Constraints in Federated Learning: Large Models Can Be Trained with only Weak Clients, *Transactions on Machine Learning Research* 2023
3. Yue Niu, Zalan Fabian, **Sunwoo Lee**, Mahdi Soltanolkotabi, and Salman Avestimehr, mL-BFGS: A Momentum-based L-BFGS for Distributed Large-Scale Neural Network Optimization, *Transactions on Machine Learning Research* 2023
4. **Sunwoo Lee**, Tuo Zhang, and Salman Avestimehr, Layer-wise Adaptive Model Aggregation for Scalable Federated Learning, *AAAI Conference on Artificial Intelligence (AAAI)*, **oral presentation** (19.7%) 2023
5. Tuo Zhang, TianTian Feng, Samiul Alam, **Sunwoo Lee**, Mi Zhang, Shrikanth S. Narayanan, and Salman Avestimehr, FedAudio: A Federated Learning Benchmark for Audio Tasks, *IEEE International Conference on Acoustics, Speech, and Signal Processing (ICASSP)* 2023
6. **Sunwoo Lee**, Chaoyang He, and Salman Avestimehr, Achieving Small-Batch Accuracy with Large-Batch Scalability via Hessian-Aware Learning Rate Adjustment. *Elsevier Neural Networks*, **158**, 1-14 2023
7. **Sunwoo Lee**, Jaeyong Jeon, and Hyungwoo Lee, Probing Oxygen Vacancy Distribution in Oxide Heterostructures by Deep Learning-based Spectral Analysis of Current Noise. *Applied Surface Science*, p154599 2022
8. **Sunwoo Lee**, Jaeyoung Jeon, Kitae Eom, Chaehwa Jeong, Yongsoo Yang, Ji-Yong Park, Chang Beom Eom, and Hyungwoo Lee, Variance-aware Weight Quantization of Multi-level Resistive Switching Devices based on Pt/LaAlO₃/SrTiO₃ Heterostructures. *Scientific Reports*, **12**, 1-10 2022
9. Kewei Wang, **Sunwoo Lee**, Jan Balewski, Alex Sim, Peter Nugent, Ankit Agrawal, Alok Choudhary, Kesheng Wu, and Wei-keng Liao, Using Multi-resolution Data to Accelerate Neural Network Training in Scientific Applications. *International Symposium on Cluster, Cloud and Internet Computing (CCGrid)* 2022
10. **Sunwoo Lee**, Qiao Kang, Reda Al-Bahrani, Ankit Agrawal, Alok Choudhary, and Wei-keng Liao, Improving Scalability of Parallel CNN Training by Adaptively Adjusting Parameter Update Frequency. *Journal of Distributed and Parallel Computing*, **159**, 10-23 2022
11. **Sunwoo Lee**, Kai-yuan Hou, Kewei Wang, Saba Sehrish, Marc Paterno, James Kowalkowski, Quincey Koziol, Ross Robert, Ankit Agrawal, Alok Choudhary, and Wei-keng Liao, A Case Study on Parallel HDF5 Dataset Concatenation for High-Energy Physics Data Analysis. *Parallel Computing*, **110**, 102877 2022
12. Kai-yuan Hou, Qiao Kang, **Sunwoo Lee**, Ankit Agrawal, Alok Choudhary, and Wei-keng Liao, Supporting Data Compression in PnetCDF, *International Conference on BigData* (19.9%) 2021

13. **Sunwoo Lee**, Qiao Kang, Kewei Wang, Jan Balewski, Alex Sim, Kesheng Wu, Ankit Agrawal, Alok Choudhary, Peter Nugent, and Wei-keng Liao, Asynchronous I/O Strategy for Large-Scale Deep Learning Applications. *International Conference on High-Performance Computing, Data, and Analytics (HiPC)* (22.9%) 2021
14. Reda Al-bahrani, Dipendra Jha, Qiao Kang, **Sunwoo Lee**, Zijiang Yang, Wei-keng Liao, Ankit Agrawal, and Alok Choudhary, SIGRNN: Synthetic minority Instances Generation in imbalanced datasets using a Recurrent Neural Network. *International Conference on Pattern Recognition Applications and Methods* 2021
15. **Sunwoo Lee**, Qiao Kang, Ankit Agrawal, Alok Choudhary, and Wei-keng Liao, Communication-Efficient Local SGD for Scalable Deep Learning. *International Conference on Big Data* (15.7%) 2020
16. Sandeep Madireddy, Ji Hwan Park, **Sunwoo Lee**, Prasanna Balaprakash, Shinjae Yoo, Wei-keng Liao, Cory Hauck, M. Paul Laiu, and Richard Archibald, In Situ Compression Artifact Removal in Scientific Data Using Deep Transfer Learning and Experience Replay. *Machine Learning: Science and Technology*, **2**, 025010 2020
17. Qiao Kang, **Sunwoo Lee**, Ankit Agrawal, Alok Choudhary, and Wei-keng Liao, Improving All-to-many Personalized Communication in MPI I/O. *International Conference for High Performance Computing, Networking, Storage, and Analysis (SC)* 2020
18. Qiao Kang, **Sunwoo Lee**, Kai-yuan Hou, Robert Ross, Ankit Agrawal, Alok Choudhary, and Wei-keng Liao, Improving MPI Collective I/O for High Volume Non-contiguous Requests with Intra-node Aggregation. *IEEE Transactions on Parallel and Distributed Systems*, **31**, 11, 2682-2695 2020
19. Qiao Kang, Alex Sim, Peter Nugent, **Sunwoo Lee**, Wei-keng Liao, Ankit Agrawal, Alok Choudhary, and Kesheng Wu. Predicting Resource Requirement in Intermediate Palomar Transient Factory Workflow. *International Symposium on Cluster, Cloud and Internet Computing (CCGrid)*, 2020 2020
20. **Sunwoo Lee**, Qiao Kang, Sandeep Madireddy, Prasanna Balaprakash, Ankit Agrawal, Alok Choudhary, Richard Archibald, and Wei-keng Liao. Improving Scalability of Parallel CNN Training by Adjusting Mini-Batch Size at Run-Time. *International Conference on Big Data* (18.7%) 2019
21. **Sunwoo Lee**, Ankit Agrawal, Prasanna Balaprakash, Alok Choudhary, and Wei-keng Liao. Communication-Efficient Parallelization Strategy for Deep Convolutional Neural Network Training. *Machine Learning in High-Performance Computing Environments (MLHPC)* 2018
22. **Sunwoo Lee**, Dipendra Jha, Ankit Agrawal, Alok Choudhary, and Wei-keng Liao. Parallel Deep Convolutional Neural Network Training by Exploiting the Overlapping of Computation and Communication. *International Conference on High-Performance Computing, Data, and Analytics (HiPC)* (22.8%) 2017

23. **Sunwoo Lee**, Wei-keng Liao, Ankit Agrawal, Nikos Hardavellas, and Alok Choudhary. 2016
Evaluation of K-Means Data Clustering Algorithm on Intel Xeon Phi. *International Conference on Big Data*
24. Diana Palsetia, William Hendrix, **Sunwoo Lee**, Ankit Agrawal, Wei-keng Liao, and Alok Choudhary. 2016
Parallel Community Detection Algorithm Using a Data Partitioning Strategy with Pairwise Subdomain Duplication. *International Conference on High Performance Computing (ISC)*
25. **Sunwoo Lee**, Byung Kwan Jung, Minsoo Ryu, Seungwon Lee, Extending 2009
Component-based Approaches for Multi-threaded Design of Multiprocessor Embedded Software. *International Symposium on Object/Component/Service-Oriented Real-Time Distributed Computing*

Pre-prints

1. Zhenheng Tang, Xiaowen Chu, Ryan Yide Ran, **Sunwoo Lee**, Shaohuai Shi, Yonggang Zhang, Yuxin Wang, Alex Qiaozhong Liang, Salman Avestimehr, Chaoyang He, FedML Parrot: A Scalable Federated Learning System via Heterogeneity-Aware Scheduling on Sequential and Hierarchical Training. *arXiv* 2023.
2. Yue Niu, Saurav Prakash, Souvik Kundu, **Sunwoo Lee**, Salman Avestimehr, Federated Learning of Large Models at the Edge via Principal Sub-Model Training. *FL-NeurIPS* 2022.
3. Chaoyang He, Zhengyu Yang, Erum Mushtaq, **Sunwoo Lee**, Mahdi Soltanolkotabi, Salman Avestimehr, SSFL: Tackling Label Deficiency in Federated Learning via Personalized Self-Supervision. *arXiv* 2021

Invited Talks

- Department of Physics at Ajou University, South Korea: System-Aware Large-Scale Neural Network Training and its Applications, 3/20/2024
- Department of Computer Engineering at Gachon University, South Korea: Scalable Federated Learning Strategies on Real-world Edge Computing Environments, 11/18/2022
- Department of Electrical Engineering at Hanyang University, South Korea: Partial Model Training Strategies in Federated Learning, 10/04/2022
- U.S. Department of Energy, SciDAC, RAPIDS Institute, Tech Talk: Asynchronous I/O Strategy for Large-Scale Deep Learning Applications, 12/01/2021
- HDF5 User Group Meeting: A Case Study on Parallel HDF5 Dataset Concatenation for Scientific Data Analysis, 10/21/2021
- U.S. Department of Energy, SciDAC, RAPIDS Institute, Tech Talk: Communication-Efficient Local SGD for Scalable Deep Learning, 7/7/2021

Service

- Program Committee of Association for the Artificial Intelligence (AAAI) 2021, 2022, 2023, 2024
- Program Committee of NeurIPS 2023, 2024
- Program Committee of KDD 2024
- Reviewer in IEEE Transactions on Parallel and Distributed Computing 2023
- Reviewer in IEEE Transactions on Mobile Computing 2023
- Program Committee of Federated Learning for Natural Language Processing (FL4NLP) workshop held in conjunction with Association for Computational Linguistics (ACL) 2022
- Program Committee of International Conference on Learning Representations (ICLR) 2021, 2022
- Program Committee of International Conference on Machine Learning (ICML) 2021, 2022, 2023
- Program Committee of International Conference on Artificial Intelligence and Statistics (AISTATS) 2021, 2022

Skills and Qualifications

Programming Language

C/C++, Python

Deep Learning Software Frameworks

TensorFlow, PyTorch, Caffe

Parallelization Libraries

MPI, OpenMP

I/O Libraries

MPI-I/O (ROMIO), HDF5, NetCDF